

Glowdays

A private skincare diary that runs one pre-registered trial at a time. This document assesses whether it should be built, what stands in the way, and in what order to proceed.

Prepared **30 July 2026**

Product **Glowdays** (formerly SkinSignal) · measurement provider **Perfect Corp YouCam Skin Analysis v2.1**

Prototype **83 screens**, Flowstep

VERDICT

The idea is sound but mis-sequenced. The differentiating method is real and validated by competitors' own workarounds. One risk could end it before product-market fit: storing facial images makes you a biometric-privacy defendant in three US states.

How to read this

Sections 1–5 are findings. Section 6 is the plan, with explicit kill criteria at each stage. If you read only one thing, read Stage 0: it costs two weeks and almost nothing, and it determines whether the product is possible at all.

Market figures are cited to their sources and, where reports disagree materially, the disagreement is reported rather than averaged. Third-party content is summarised and rephrased for licensing compliance. Nothing here is legal advice; the privacy section identifies exposure that requires review by qualified counsel.

Market size, and why the published numbers are not a plan

The space is real and growing. The sizing literature is not usable for decisions.

For overlapping definitions of AI skin analysis in 2026, published estimates span more than a factor of seven:

SOURCE	2026 ESTIMATE	STATED CAGR	HORIZON
360iResearch	USD 0.83bn	15.9%	2.02bn by 2032
Mordor Intelligence	USD 1.02bn	19.3%	2.45bn by 2031
Precedence Research	USD 2.09bn	16.2%	8.02bn by 2035
Coherent (via EIN)	USD 2.13bn	16.8%	6.30bn by 2033
Towards Healthcare	USD 5.84bn	18.7%	27.41bn by 2035

Finding

A 7× spread across the same nominal market in the same year means these are marketing artefacts, not measurements. Do not put any of these numbers in a deck or a grant application as though they were facts. Cite the *direction* — consistent mid-to-high-teens growth — and nothing more.

The uncomfortable part

What the reports do agree on is *who is buying*: beauty brands, retailers, clinics and med-spas embedding analysis into their own funnels. Perfect Corp positions the technology as something brands and clinics deploy, citing deployments with major dermatological skincare names.

The revenue in this category today is **B2B tooling sold to brands, not consumer subscriptions**. Glowdays is a consumer product, which is the thinner half of a growing market. That does not make it unviable, but it means the market-growth story is not evidence for *your* business model. It is evidence for your supplier's.

Sources: report pages listed above; [Perfect Corp AI Skin Analysis](#). Content rephrased for licensing compliance.

Competitive landscape

Three clusters. You compete with one of them, and your supplier sits in another.

CLUSTER	PLAYERS	WHAT THEY DO	RELATION TO GLOWDAYS
DIRECT Consumer trackers	Skin Bliss, TroveSkin	Face scan, routine building, product scoring, progress photos, recommendations	Head-on. Skin Bliss cites a 50,000- respondent satisfaction survey, so it has real scale and distribution.
ADJACENT Medical monitoring	SkinVision, Miiskin, mole trackers	Lesion and melanoma surveillance, some regulated as medical devices	Not competitors, but they shape what users expect a "skin app" to claim, and they raise the regulatory baseline in users' minds.
SUPPLIER B2B white- label	Perfect Corp, Haut.ai, Reveue	Sell analysis APIs and turnkey experiences to brands	Perfect Corp is simultaneously your measurement supplier and a potential competitor. It publishes guides on building profitable skin-analysis apps on its own API.

Structural risk: no moat at the analysis layer

Anyone can buy the same scores you buy, from the same vendor, at the same quality. Your defensibility is entirely in the *method* layer — pre-registration, single-variable isolation, confidence gating, and refusal to show unsupportable comparisons. Protect that, and accept that the measurement itself is a commodity.

Sources: [Skin Bliss](#); [Perfect Corp app roundup](#); [Perfect Corp build guide](#). Rephrased for compliance.

The strongest validation is accidental

Competitors' own guidance reinvents your product thesis as a workaround.

Independent guidance for using TroveSkin advises

treating the app as a progress journal rather than a diagnostic tool; photographing in the same light and angle each time; changing one habit at a time; and reading week-to-week trends instead of day-to-day judgements.

Paraphrased from [apkafe TroveSkin guide](#).

That is Glowdays, described as a set of manual corrections a user must apply to make somebody else's app trustworthy. Four separate elements of your design — journal-not-diagnostic framing, capture consistency, single-variable isolation, and reading over weeks — appear in third-party advice about a competitor.

Separately, Skin Bliss acknowledges that a blurry photo or glare can drop ingredients from a scan and shift a score meaningfully. An incumbent conceding that its scores move with capture conditions is precisely the problem your confidence model exists to solve.

What this does and does not prove

- **Proves:** the problem is real, users encounter it, and the workaround is already being taught informally.
- **Does not prove:** that users will pay for rigour. Methodology is harder to market than a score. "See your skin score" converts; "see whether your comparison is trustworthy" is a harder sell that appeals to a narrower, more patient audience.

Source: [Skin Bliss FAQ](#). Rephrased for compliance.

The existential risk: biometric privacy

This is the section that should change your launch plan.

Glowdays is a wellness app that collects facial images, derives inferences about the user's body, and transmits those images to a third-party processor. That is close to the paradigm case for the current wave of US biometric-privacy litigation.

REGIME	WHY IT REACHES YOU	EXPOSURE
Washington My Health My Data Act	Consumer health data definitions reach facial imagery from which an identifier template can be extracted. Being outside HIPAA is expressly not a safe harbour for wellness apps.	PRIVATE RIGHT OF ACTION via the state Consumer Protection Act. Described by practitioners as "BIPA 2.0".
Illinois BIPA	Facial geometry collection without compliant written consent and a published retention schedule.	STATUTORY DAMAGES PER VIOLATION , the most litigated biometric regime in the US.
Texas CUBI	Biometric identifiers captured for commercial purpose.	Attorney-General enforcement, substantial per-violation penalties.
GDPR Art. 9	Biometric and health data are special category; requires explicit consent and DPIA.	Relevant the moment you accept an EU user.

Class actions under the Washington statute have already been filed, including a matter premised on SDKs embedded in mobile apps collecting consumer health data. The litigation wave predicted when the Act passed has begun.

Why this is existential rather than merely serious

A single certified class action would cost more than this product will earn in several years. For a bootstrapped consumer app, biometric exposure is not a compliance line item — it is a survival question, and it must be resolved *before* user acquisition rather than after.

Mitigations, in order of cost-effectiveness

1. **Geofence at launch.** Exclude Washington, Illinois and Texas until consent flows are counsel-reviewed. Nearly free, buys months.
2. **Say "biometric".** Your consent sheet currently says "facial image". Compliant consent typically must name biometric collection explicitly, state purpose, name the processor, and publish a retention and destruction schedule.

3. **Version the consent and re-consent on change.** Already designed; make sure the stored record includes the version and a hash of the text shown.
4. **Get the retention term contractually.** The 24-hour Perfect Corp retention window appears on five user-facing screens as a factual promise, sourced from documentation rather than from your contract. Confirm it in writing before launch.
5. **Consider on-device-only originals.** Upload for the analysis call, keep no server copy of the source image. This materially reduces exposure at the cost of cross-device history. Worth modelling as a product decision, not just a security one.

Sources: [IAPP on MHMDA](#); [first MHMDA lawsuit](#); [SDK-based class action](#); [MHMDA and wellness apps](#); [Steptoe on private right of action](#). Summarised; not legal advice.

Two data problems that constrain the product

5.1 The vendor's accuracy claim is internally inconsistent

Across Perfect Corp's own pages the same capability is described as:

CLAIM	WHERE
95% test-retest reliability, verified by medical study	Skin Analysis API page
95% clinical accuracy, dermatologist-grade	AI Skin Diagnostic page
80% correlation with dermatologist evaluations	Skin Analysis tool page

Test-retest reliability, clinical accuracy and correlation with expert rating are three different quantities. A product whose entire positioning is honesty about measurement cannot repeat any of them. You need your own number.

Gate zero

Two HD scans of the same face in one sitting, same light, repeated across roughly ten people. This yields σ , the test-retest spread per metric, which sets the smallest movement you are entitled to call a finding. **If σ exceeds roughly 8 raw points on hydration, the app cannot detect the effects it promises and the measurement claim must be dropped.**

5.2 Pooled evidence cannot be a launch feature

Your own admission threshold is 30 single-variable, high-detail, medium-confidence-or-better trials per product. The arithmetic of populating that is unforgiving:

6,000	~30,000	1
CLEAN TRIALS TO POPULATE 200 PRODUCTS	COMMITTED USERS, AT 20% COMPLETION	SCREEN USERS WILL ACTUALLY SEE AT LAUNCH

Which means **the withheld-evidence state is your launch experience**. It is a good screen and an honest one, but it cannot carry acquisition. Sequence pooling late and market something else first.

5.3 Unit economics: unresolved

Perfect Corp does not publish per-call API pricing; it is quote-based. This is the single largest open number in the business case. At two check-ins per user per month, plus storage of HD result archives and masks, cost per retained user determines whether any subscription price works. **Obtain a written quote before further build.**

Action plan

Four stages, each with an explicit kill criterion. Do not begin a stage before the previous gate passes.

STAGE 0 · TWO WEEKS · BEFORE ANY FURTHER BUILD

Establish that measurement is possible

- Obtain written API pricing and the contractual image-retention term from Perfect Corp.
- Run the test-retest study: 10 participants, two HD scans in one sitting, consistent light. Compute σ per metric.
- Re-derive the minimum detectable movement from σ and update every threshold in the product accordingly.

Kill criterion

σ on the headline metric exceeds ~8 raw points → stop. Ship a pre-registered journal without scores, or abandon.

STAGE 1 · FOUR WEEKS · CONCIERGE VALIDATION

Prove people can produce comparable pairs

- Twenty recruited users, one pre-registered trial each, eight weeks. No pooling, no shelf, no calendar.
- Run the five-task usability script from the retention strategy, against its stated pass criteria: 4 of 5 start a scan unaided; 4 of 5 explain "not enough evidence" correctly; nobody reads a raw score as a diagnosis; 4 of 5 find deletion within 60 seconds.
- Geofenced. Counsel-reviewed consent in place before the first upload.

Kill criterion

Fewer than half produce two comparable check-ins → the capture-consistency problem is unsolved, and UI will not fix it.

STAGE 2 · CLOSED BETA

Test whether the loop retains without coercion

- Add shelf, calendar, rhythm, drafts, export. Still geofenced.
- Instrument the north-star metric already defined: *meaningful comparison rate* — new trials producing two successful scans within the chosen cadence whose owner opens the comparison.

- Treat the harm guardrails as hard stops on any retention experiment: notification opt-out rate, notification disable rate, track and account deletion rate, support complaints.

Kill criterion

Retention rises while deletion or opt-out also rises → per your own strategy document, that is failure, not growth.

STAGE 3 · AFTER ~10,000 ACTIVE USERS

Turn on pooled evidence

- Publish the pooling methodology before publishing any pooled number.
- Enable only for products past the 30-trial threshold. Everything else keeps the withheld state.
- Never rank, never recommend, never accept sponsorship. Continuation rate and distribution only.

Monetisation

Subscription is the only model compatible with the product's own non-goals. Affiliate or commission revenue would corrupt pooled evidence, which is the sole durable asset. Expect low conversion and high retention among converts — a small, durable business rather than a venture-scale one. Decide explicitly which you are building, because the answer changes Stage 2's success threshold.

What would change the assessment

If σ comes back small and Stage 1 shows users can reliably produce comparable pairs, this is a defensible product with a genuine method advantage. If σ is large, the honest move is to drop the measurement claim entirely and ship the pre-registered journal without scores — still useful to people, and a much smaller business.